



OTTO VON GUERICKE
UNIVERSITÄT
MAGDEBURG

INF

FAKULTÄT FÜR
INFORMATIK

SIMD Acceleration for Index Structures

Marten Wallewein-Eising

Otto von Guericke University, Magdeburg

January 23, 2018

Agenda

Motivation

Excursion: B^+ -Tree

SIMD Style Processing

Adapted Tree structures

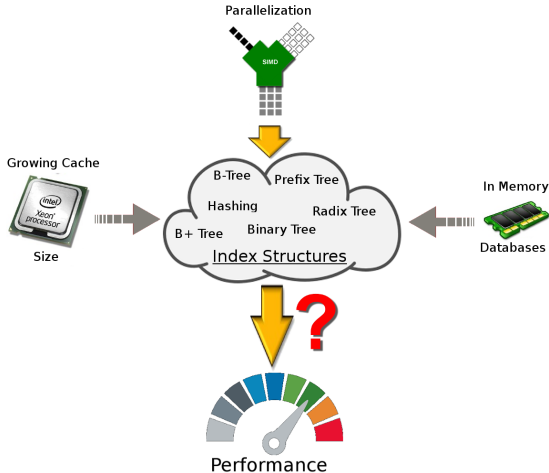
Seg-Tree/Trie

FAST

Evaluation

Conclusion

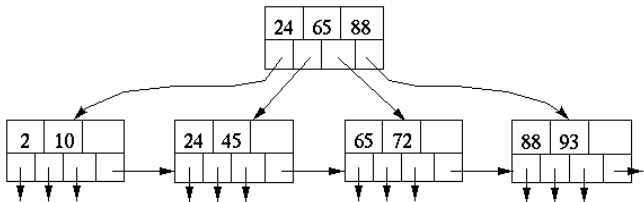
Motivation



Excursion: B⁺-Tree

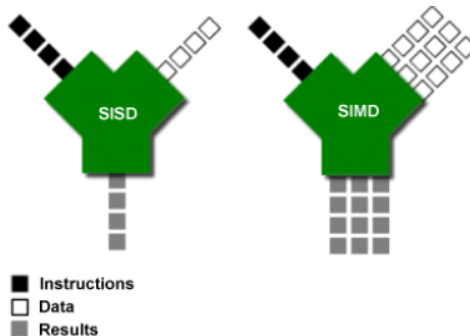


B⁺-Tree



- N-ary tree with large number of children per node
- Only leaf nodes contain values
- Leaf nodes often linked for range based scans

Single Instruction Multiple Data

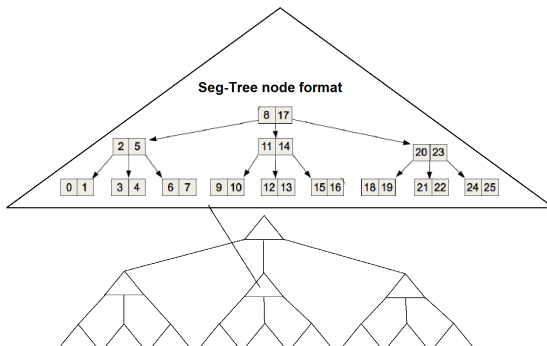


- `__m128i _mm_cmpgt_epi32 (__m128i a, __m128i b)`
Compares 4 signed 32-bit integers in a and 4 signed 32-bit integers in b for greater-than.

Adapted Tree structures

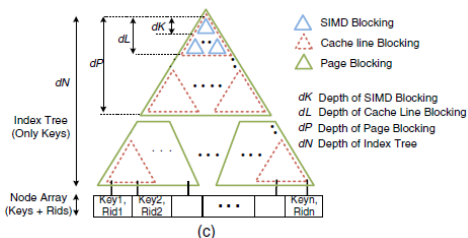
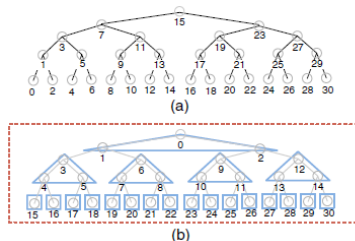
- Seg-Tree/Trie
- FAST: Fast Architecture Sensitive Tree

Seg-Tree/Trie



- Each node is a k-ary search tree
- Each node is linearised to use k-ary search
- $k = \frac{|SIMD|}{|Key|}$, k keys are compared in parallel

Fast Architecture Sensitive Tree



- Based on binary tree
- Hierarchical blocking: SIMD, cache line and page blocks
- Efficient cache line and page usage

Evaluation

Important criteria for performance increase:

- Horizontal vectorization
- Minimized key size
- Adapted node sizes and types
- Decreased branch misses
- Full use of cache line using blocking and alignment
- Usage of compression
- Adapt search algorithm for linearised nodes

Evaluation

Implementation of the considered performance criteria and their impact:

Criterion	Seg-Tree/Trie	FAST	ART	VAST	Impact
Horizontal vectorization	x	x	x	x	high
Minimized key size	o	-	x	x	high
Adapted node sizes and types	-	x	-	x	low
Decreased branch misses	-	x	-	x	medium
Full use of cache line using blocking and alignment	-	x	-	x	medium
Usage of Compression	o	-	x	x	medium
Adapt search algorithm for linearised nodes	x	-	-	-	low

Legend: x: implements the issue, o: partially implements the issue, -: not implements the issue

Conclusion

How to adapt index structures to modern database systems:

- Compare as many keys as possible in parallel with SIMD
 - Direct performance increase up to a multiple
- Efficient usage of cache line
- Decrease branch misses
- Use compression or adapted search algorithms

Sources

- <http://infolab.stanford.edu/~nsample/cs245/handouts/hw2sol/sol2.html>
- https://en.wikipedia.org/wiki/Radix_tree
- <https://www.clker.com/clipart-bosque.html>

Thank you for your attention!