

SIMD Acceleration for Index Structures: A Survey

Beyond Databases Architectures and Structures

Marten Wallewein-Eising, Gunter Saake, David Broneske

University of Magdeburg, Germany

September 18, 2018



Agenda

Motivation

SIMD Style Processing

Surveyed Main-Memory Index Structures

Elf

Seg-Tree/Trie

FAST

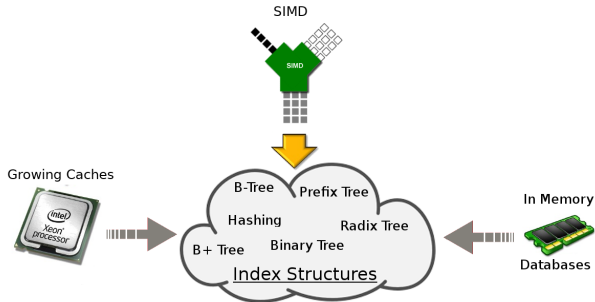
VAST

ART

Qualitative Comparison

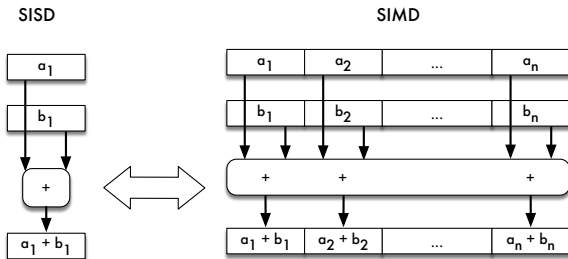
Conclusion

Motivation



- What are state-of-the-art main-memory index structures?
- Which optimizations do they have in common?

Single Instruction Multiple Data



- `__m128i _mm_add_epi32 (__m128i a, __m128i b)` Adds 4 signed 32-bit integers in **a** to 4 signed 32-bit integers in **b**.

Surveyed Main-Memory Index Structures

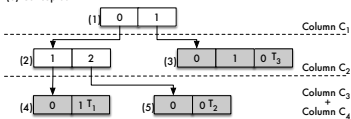
Index Structure	Based on	Reference
Elf	Prefix Tree	[Broneske et al. ICDE'17]
Seg-Tree/Trie	CSB-Tree	[Zeuch et al. EDBT'14]
FAST: Fast Architecture Sensitive Tree	Binary Tree	[Kim et al. SIGMOD'10]
VAST: Vector-Advanced and Compressed Structure Tree	Binary Tree	[Yamamuro et al. EDBT'12]
ART: Adaptive Radix Tree	Radix Tree	[Leis et al. ICDE'13]

Elf

(a) Input Table

C ₁	C ₂	C ₃	C ₄	Ref
0	1	0	1	T ₁
0	2	0	0	T ₂
1	0	1	0	T ₃

(b) Conceptual Elf



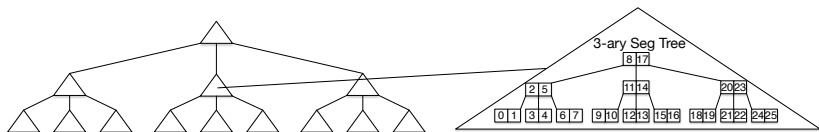
(c) Optimized Memory Layout for Elf

	0	1	2	3	4	5	6	7	8	9
Elf[00]	⁽¹⁾ [02]	⁽²⁾ [-12]	1	[-6]	-2	[-9]	0	1	T ₁	⁽⁵⁾ 0
Elf[10]	0	T ₂	⁽³⁾ 0	1	0	T ₃				
Elf[20]										

- Multi-dimensional index structure for column-wise storage
- Prefix-redundancy elimination on distinct column values
- Linearisation for optimized memory layout

[Broneske et al. ICDE'17]

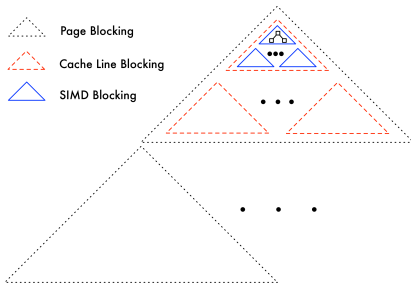
Seg-Tree/Trie



- Each node is a k-ary search tree
- Each node is linearised to use k-ary search
- $k = \frac{|SIMD|}{|Key|} + 1$ partitions, $k - 1$ separator keys are compared in parallel

[Zeuch et al. EDBT'14]

Fast Architecture Sensitive Tree



- Based on binary tree
- Hierarchical blocking: page, cache line and SIMD blocks
- Efficient register, cache line and page usage

[Kim et al. SIGMOD'10]

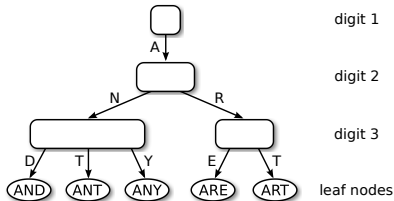


Vector-Advanced and Compressed Structure Tree

- Extension of FAST
- Decrease branch misses avoiding conditional branches
- Uses key compression
 - Lossy compression for inner nodes
 - Lossfree compression leaf nodes
- Decompression and error correction of lossy compression has less impact compared to the performance increase with SIMD

[Yamamuro et al. EDBT'12]

Adaptive Radix Tree



- Uses different node types with different number of keys and children
- Due to overfill or underfill of nodes, the node type is changed
- Reduced space consumption due to lazy expansion and path compression

[Leis et al. ICDE'13]

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high
Minimized key size	o	-	x	x	-	high

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high
Minimized key size	o	-	x	x	-	high
Adapted node sizes / types	-	x	-	x	-	low

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high
Minimized key size	o	-	x	x	-	high
Adapted node sizes / types	-	x	-	x	-	low
Decreased branch misses	-	x	-	x	-	medium

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high
Minimized key size	o	-	x	x	-	high
Adapted node sizes / types	-	x	-	x	-	low
Decreased branch misses	-	x	-	x	-	medium
Exploit cache lines using blocking and alignment	-	x	-	x	x	medium

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high
Minimized key size	o	-	x	x	-	high
Adapted node sizes / types	-	x	-	x	-	low
Decreased branch misses	-	x	-	x	-	medium
Exploit cache lines using blocking and alignment	-	x	-	x	x	medium
Use of Compression	o	-	x	x	x	medium

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact

Qualitative Comparison

Criterion	Seg-Tree/ Trie	FAST	ART	VAST	Elf	Impact
Horizontal vectorization	x	x	x	x	-	high
Minimized key size	o	-	x	x	-	high
Adapted node sizes / types	-	x	-	x	-	low
Decreased branch misses	-	x	-	x	-	medium
Exploit cache lines using blocking and alignment	-	x	-	x	x	medium
Usage of Compression	o	-	x	x	x	medium
Adapt search algorithm for linearized nodes	x	-	-	-	x	low

Legend: x = implements the issue; o = partially implements the issue;
- = does not implement the issue

Table: Implementation of the considered performance criteria and their impact



Conclusion

How to adapt index structures to modern database systems:

- Compare as many keys as possible in parallel with SIMD
 - Direct performance increase by a factor of X (where X keys fit a SIMD register)
- Efficient usage of cache line
- Decrease branch misses
- Use compression or/and adapted search algorithms



Conclusion

How to adapt index structures to modern database systems:

- Compare as many keys as possible in parallel with SIMD
 - Direct performance increase by a factor of X (where X keys fit a SIMD register)
- Efficient usage of cache line
- Decrease branch misses
- Use compression or/and adapted search algorithms

Thank you for your attention!

References I

-  Broneske, D., Köppen, V., Saake, G., and Schäler, M. (2017). Accelerating multi-column selection predicates in main-memory - the Elf approach.
In Proceedings of the International Conference on Data Engineering (ICDE), pages 647–658. IEEE.
-  Kim, C., Chhugani, J., Satish, N., Sedlar, E., Nguyen, A. D., Kaldewey, T., Lee, V. W., Brandt, S. A., and Dubey, P. (2010). Fast: Fast architecture sensitive tree search on modern CPUs and GPUs.
In Proceedings of the International Conference on Management of Data (SIGMOD), pages 339–350. ACM.



References II



Leis, V., Kemper, A., and Neumann, T. (2013).
The adaptive radix tree: Artful indexing for main-memory
databases.

In *Proceedings of the International Conference on Data
Engineering (ICDE)*, pages 38–49. IEEE.



Yamamuro, T., Onizuka, M., Hitaka, T., and Yamamuro, M.
(2012).

Vast-tree: A vector-advanced and compressed structure for
massive data tree traversal.

In *Proceedings of the International Conference on Extending
Database Technology (EDBT)*, pages 396–407. ACM.



References III



Zeuch, S., Huber, F., and Freytag, J.-c. (2014).
Adapting tree structures for processing with simd instructions.
In *Proceedings of the International Conference on Extending
Database Technology (EDBT)*. Citeseer.